

Task 1: Execute WordMedian , WordMean , WordStandardDeviation programs using hadoop-mapreduce-examples-2.9.0.jar file present in your AcadGild VM.

File to execute wordmedian,wordmean, wordstandarddeviation:-

```
[acadgild@localhost mapreduce]$ hadoop fs -cat /hadoopdata/hdfs/file1.txt
18/02/21 16:55:17 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Hello,
```

```
Welcome to acadgild!This program is regarding to training on Big Data Hadoop and Spark training.
```

```
Thank you all!
```

Executing word mean

Executed jar files for word mean using the command as highlighted in the below screenshot and wordmean of the file "**File1.txt**" is **5.26**.

```
acadgild@localhost:~/install/hadoop/hadoop-2.6.5/share/hado
File Edit View Search Terminal Help
Usage: wordmean <in> <out>
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost mapreduce]$ hadoop jar hadoop-mapreduce-examples-2.6.5.jar w
ordmean /hadoopdata/hdfs/file1.txt /hadoopdata/hdfs-output
18/02/21 17:19:14 WARN util.NativeCodeLoader: Unable to load native-hadoop libra
ry for your platform... using builtin-java classes where applicable
18/02/21 17:19:20 INFO client.RMPProxy: Connecting to ResourceManager at localhos
t/127.0.0.1:8032
18/02/21 17:19:27 INFO input.FileInputFormat: Total input paths to process : 1
18/02/21 17:19:29 INFO mapreduce.JobSubmitter: number of splits:1
18/02/21 17:19:30 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_15
19210873612_0004
18/02/21 17:19:32 INFO impl.YarnClientImpl: Submitted application application_15
19210873612_0004
18/02/21 17:19:33 INFO mapreduce.Job: The url to track the job: http://localhost
:8088/proxy/application_1519210873612_0004/
18/02/21 17:19:33 INFO mapreduce.Job: Running job: job_1519210873612_0004
18/02/21 17:20:28 INFO mapreduce.Job: Job job_1519210873612_0004 running in uber
mode : false
18/02/21 17:20:28 INFO mapreduce.Job: map 0% reduce 0%
18/02/21 17:21:12 INFO mapreduce.Job: map 100% reduce 0%
18/02/21 17:21:53 INFO mapreduce.Job: map 100% reduce 100%
18/02/21 17:21:57 INFO mapreduce.Job: Job job_1519210873612_0004 completed succe
ssfully
18/02/21 17:21:58 INFO mapreduce.Job: Counters: 49
    File System Counters
        FILE: Number of bytes read=39
        FILE: Number of bytes written=215445
        FILE: Number of read operations=0
        FILE: Number of large read operations=0
        FILE: Number of write operations=0
        HDFS: Number of bytes read=233
        HDFS: Number of bytes written=20
        HDFS: Number of read operations=6
        HDFS: Number of large read operations=0
        HDFS: Number of write operations=2
    Job Counters
        Launched map tasks=1
        Launched reduce tasks=1
```

```
acadgild@localhost:~/install/hadoop/hadoop-2.6.5/share/had
File Edit View Search Terminal Help
Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=39561
    Total time spent by all reduces in occupied slots (ms)=37722
    Total time spent by all map tasks (ms)=39561
    Total time spent by all reduce tasks (ms)=37722
    Total vcore-milliseconds taken by all map tasks=39561
    Total vcore-milliseconds taken by all reduce tasks=37722
    Total megabyte-milliseconds taken by all map tasks=40510464
    Total megabyte-milliseconds taken by all reduce tasks=38627328
Map-Reduce Framework
    Map input records=5
    Map output records=38
    Map output bytes=551
    Map output materialized bytes=39
    Input split bytes=112
    Combine input records=38
    Combine output records=2
    Reduce input groups=2
    Reduce shuffle bytes=39
    Reduce input records=2
    Reduce output records=2
    Spilled Records=4
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=944
    CPU time spent (ms)=8660
    Physical memory (bytes) snapshot=292208640
    Virtual memory (bytes) snapshot=4120752128
    Total committed heap usage (bytes)=170004480
Shuffle Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0
File Input Format Counters
    Bytes Read=121
File Output Format Counters
    Bytes Written=20
The mean is: 5.2631578947368425
```

Executing WordMedian:

Executed jar files for word median using the command as highlighted in the below screenshot and wordmean of the file "**File1.txt**" is 5.

```
acadgild@localhost:~/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce
File Edit View Search Terminal Help
asses where applicable
You have new mail in /var/spool/mail/acadgild
[acadgild@localhost mapreduce]$ hadoop jar hadoop-mapreduce-examples-2.6.5.jar wordmedian /hadoopdata/hdfs/file1.txt /hadoopdata/hdfs/results-output
18/02/21 17:26:30 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/02/21 17:26:36 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/02/21 17:26:41 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your application with ToolRunner to remedy this.
18/02/21 17:26:44 INFO input.FileInputFormat: Total input paths to process : 1
18/02/21 17:26:45 INFO mapreduce.JobSubmitter: number of splits:1
18/02/21 17:26:46 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1519210873612_0005
18/02/21 17:26:48 INFO impl.YarnClientImpl: Submitted application application_1519210873612_0005
18/02/21 17:26:49 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1519210873612_0005/
18/02/21 17:26:49 INFO mapreduce.Job: Running job: job_1519210873612_0005
18/02/21 17:27:42 INFO mapreduce.Job: Job job_1519210873612_0005 running in uber mode : false
18/02/21 17:27:42 INFO mapreduce.Job: map 0% reduce 0%
18/02/21 17:28:27 INFO mapreduce.Job: map 100% reduce 0%
18/02/21 17:29:10 INFO mapreduce.Job: map 100% reduce 100%
18/02/21 17:29:12 INFO mapreduce.Job: Job job_1519210873612_0005 completed successfully
18/02/21 17:29:14 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=96
    FILE: Number of bytes written=215311
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=233
    HDFS: Number of bytes written=37
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=38992
    Total time spent by all reduces in occupied slots (ms)=38903
    Total time spent by all map tasks (ms)=38992
```



```
acadgild@localhost:~/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce
File Edit View Search Terminal Help
Total vcore-milliseconds taken by all map tasks=38992
Total vcore-milliseconds taken by all reduce tasks=38903
Total megabyte-milliseconds taken by all map tasks=39927808
Total megabyte-milliseconds taken by all reduce tasks=39836672
Map-Reduce Framework
  Map input records=5
  Map output records=19
  Map output bytes=152
  Map output materialized bytes=96
  Input split bytes=112
  Combine input records=19
  Combine output records=9
  Reduce input groups=9
  Reduce shuffle bytes=96
  Reduce input records=9
  Reduce output records=9
  Spilled Records=18
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=694
  CPU time spent (ms)=8760
  Physical memory (bytes) snapshot=293756928
  Virtual memory (bytes) snapshot=4117938176
  Total committed heap usage (bytes)=170004480
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=121
File Output Format Counters
  Bytes Written=37
The median is: 5
```

Executing WordStandardDeviation:

Executed jar files for word standarddeviation using the command as highlighted in the below screenshot and wordmean of the file "**File1.txt**" is **2.93**.

```
acadgild@localhost:~/install/hadoop/hadoop-2.6.5/share/hadoop/mapreduce
File Edit View Search Terminal Help
[acadgild@localhost mapreduce]$ hadoop jar hadoop-mapreduce-examples-2.6.5.jar wordstandarddeviation /hadoopdata/hdfs/file1.txt /hadoopdata/hdfs/resul-output
18/02/21 18:07:18 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
18/02/21 18:07:25 INFO client.RMProxy: Connecting to ResourceManager at localhost/127.0.0.1:8032
18/02/21 18:07:34 INFO input.FileInputFormat: Total input paths to process : 1
18/02/21 18:07:35 INFO mapreduce.JobSubmitter: number of splits:1
18/02/21 18:07:37 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1519210873612_0007
18/02/21 18:07:40 INFO impl.YarnClientImpl: Submitted application application_1519210873612_0007
18/02/21 18:07:41 INFO mapreduce.Job: The url to track the job: http://localhost:8088/proxy/application_1519210873612_0007/
18/02/21 18:07:41 INFO mapreduce.Job: Running job: job_1519210873612_0007
18/02/21 18:08:42 INFO mapreduce.Job: Job job_1519210873612_0007 running in uber mode : false
18/02/21 18:08:42 INFO mapreduce.Job: map 0% reduce 0%
18/02/21 18:09:45 INFO mapreduce.Job: map 100% reduce 0%
18/02/21 18:10:44 INFO mapreduce.Job: map 100% reduce 100%
18/02/21 18:10:51 INFO mapreduce.Job: Job job_1519210873612_0007 completed successfully
18/02/21 18:10:52 INFO mapreduce.Job: Counters: 49
  File System Counters
    FILE: Number of bytes read=56
    FILE: Number of bytes written=215651
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=233
    HDFS: Number of bytes written=31
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
  Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=57993
    Total time spent by all reduces in occupied slots (ms)=57034
    Total time spent by all map tasks (ms)=57993
    Total time spent by all reduce tasks (ms)=57034
    Total vcore-milliseconds taken by all map tasks=57993
    Total vcore-milliseconds taken by all reduce tasks=57034
    Total megabyte-milliseconds taken by all map tasks=59384832
```

```
acadgild@localhost:~/install/hadoop/hadoop-2.6.5/share/hadoop
File Edit View Search Terminal Help
Total vcore-milliseconds taken by all map tasks=57993
Total vcore-milliseconds taken by all reduce tasks=57034
Total megabyte-milliseconds taken by all map tasks=59384832
Total megabyte-milliseconds taken by all reduce tasks=58402816
Map-Reduce Framework
  Map input records=5
  Map output records=57
  Map output bytes=836
  Map output materialized bytes=56
  Input split bytes=112
  Combine input records=57
  Combine output records=3
  Reduce input groups=3
  Reduce shuffle bytes=56
  Reduce input records=3
  Reduce output records=3
  Spilled Records=6
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=976
  CPU time spent (ms)=9860
  Physical memory (bytes) snapshot=296062976
  Virtual memory (bytes) snapshot=4117938176
  Total committed heap usage (bytes)=170004480
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=121
File Output Format Counters
  Bytes Written=31
The standard deviation is: 2.935124946020775
```